


# Operating Systems

Autumn 2024

OS abstraction: process

# Textbook: "OSTEP"

Operating Systems: Three Easy Pieces (Available for free online)



## Virtualization

### Virtualizing CPU

- Processes
- Threads
- Scheduling ...

### Virtualizing Memory

- Address space
- Segmentation
- Paging ...

## Concurrency

- Threads
- Locks
- Conditional variables
- Semaphores ...

## Persistence

- File systems
- Journaling, logging
- I/O devices
- File integrity ...

# Marking scheme

- Two in-sems: 25%
- End-sem: 50%
- Lab exercises: 15%
- Mini projects: 10%

# Process $\neq$ Program

- A process is a running instance of a program
  - Program = static file (image)
  - Process = executing program = program + execution state
- Several processes may run the same program code, but each is a distinct process with its own state
- A process executes sequentially, one instruction at a time
- Two processes are said to run concurrently when instructions of one process are interleaved with the instructions of the other process

# Process states

- ready: waiting to be assigned to CPU
  - could run, but another process has the CPU

# Process states

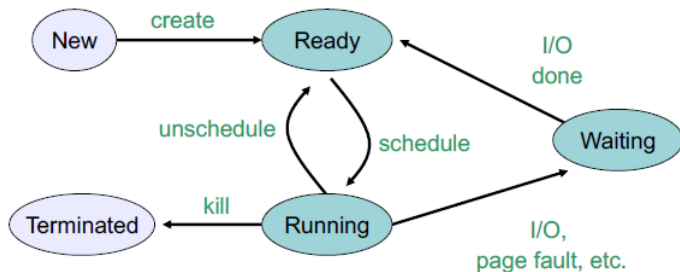
- ready: waiting to be assigned to CPU
  - could run, but another process has the CPU
- running: executing on the CPU
  - is the process that currently controls the CPU

# Process states

- **ready**: waiting to be assigned to CPU
  - could run, but another process has the CPU
- **running**: executing on the CPU
  - is the process that currently controls the CPU
- **waiting**: waiting for an event, e.g. I/O
  - cannot make progress until event happens



# Process state transitions



Question: What can cause schedule/unschedule transitions?

# Example: state transitions

New  
Ready  
Running

```
main() {  
    printf("Hello world");  
}
```

## Example: state transitions

```
main() {  
    printf("Hello world");  
}
```

New  
Ready  
Running  
  
Waiting

# Example: state transitions

```
main() {  
    printf("Hello world");  
}
```

New  
Ready  
Running

Waiting

Ready  
Running  
Terminated

# Context Switch

- Switching the CPU from one process to another is called a context switch – relatively expensive operation
- Time sharing systems may do 100 to 1000 context switches a second

# Creating a process

- One process can create other processes to do work
  - The creator is called the parent and the new process is the child
  - A parent can either wait for the child to complete, or continue in parallel

# Creating a process

- One process can create other processes to do work
  - The creator is called the parent and the new process is the child
  - A parent can either wait for the child to complete, or continue in parallel
- In Unix, the `fork()` system call is used to create child processes

# System calls

- OS procedures that perform privileged operations
  - Linux x86\_64 has ~323 system calls, numbered from 0–322
  - OS uses a `sys_call_table` to keep the syscall handlers (indexed by syscall number)
- Now process is able to perform ~300 different kinds of restricted operations, and cannot access anything it wants to



# Creating a process

- `fork()` copies variables and registers from the parent to the child
- `fork()`, when called, returns twice (to each process)
- Return value of `fork()`
  - In the parent process, `fork()` returns the process id of the child
  - In the child process, the return value is 0

# Creating a process

- `fork()` copies variables and registers from the parent to the child
- `fork()`, when called, returns twice (to each process)
- Return value of `fork()`
  - In the parent process, `fork()` returns the process id of the child
  - In the child process, the return value is 0
- The parent can wait for the child to terminate by executing the `wait` system call or continue execution

## pop quiz

```
main() {  
    int x = 0;  
    int cid = fork();  
    if (cid == 0) {  
        x = 9; printf("%d ", x);  
    }  
    else {  
        x = 10; printf("%d ", x);  
    }  
}
```



This code will print out:

- (A) 10 10
- (B) 9 9
- (C) 10 9
- (D) 9 10

- order of execution is non-deterministic
  - parent and child run concurrently
- Important: post fork, parent and child are identical but separate!
  - OS allocates and maintains separate data/state
  - control flow can diverge

# Fun stuff: process graphs

```
void fork2() {  
    printf("L0\n");  
    fork();  
    printf("L1\n");  
    fork();  
    printf("Bye\n");  
}
```



Okay, `fork()` creates a new process that is a duplicate of the process ...

What if I want to run something different?

- It does NOT create a new process
- It basically replaces the current process with a new program
- Does `exec()` ever return? If so, what does it mean?

All functions (`exec()`-family) return -1 in the case of an error. Otherwise at successful execution there is no return back to the calling program. Thus, it is redundant to check the return value; you can directly continue with the error routine

Normally, I want to start a different program without replacing the current process (type "firefox" from a terminal).

How can we do this?

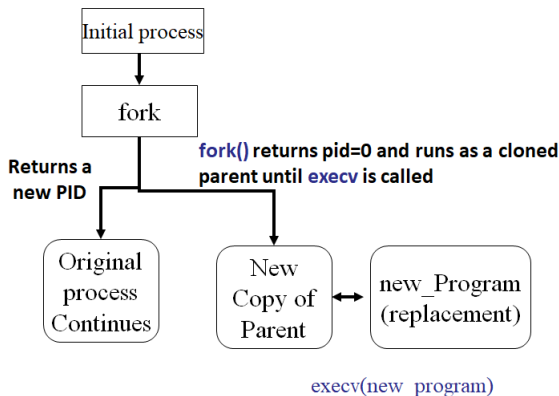


Normally, I want to start a different program without replacing the current process (type "firefox" from a terminal).

How can we do this?

Answer: call `fork()` then `exec()`

```
execv(new_program, argv[ ])
```



# The scheduling problem

- Which process should the OS run?
  - if no runnable process (i.e. no process in the ready state), run the idle task
  - if a single process runnable, run this one
  - if more than one runnable process, a scheduling decision must be taken
- Scheduler
  - code (logic) that decides which process to run as per some policy

# Coming next

- What are the basic scheduling policies?
- When do they work well?