

# Operating Systems

Autumn 2024

CPU scheduling

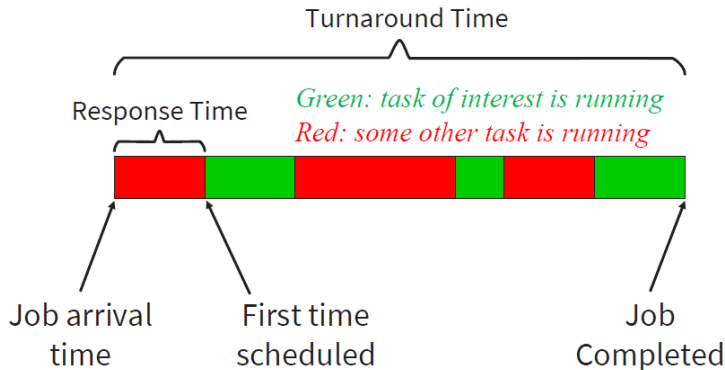
# The scheduling problem

- Which process should the OS run?
  - if no runnable process (i.e. no process in the ready state), run the idle task
  - if a single process runnable, run this one
  - if more than one runnable process, a scheduling decision must be taken
- Scheduler: code (logic) that decides which process to run as per some policy
- Today: What are the basic scheduling policies? When do they work well?

# Standard usage

- Refer to schedulable entities as jobs — could be processes, threads, etc.
- Job: a task that needs a period of CPU time
- Job arrival time
  - when the job was first submitted
- Job run time
  - Time needed to run the task without contention

# Scheduling metrics



- Execution time: sum of green periods
- Waiting time: sum of red periods
- Turnaround time: sum of both

# Performance terminology

- Turnaround time: How long?
  - User-perceived time to complete some job ( $\text{completion\_time} - \text{arrival\_time}$ )
- Response time:
  - User-perceived time before first output ( $\text{initial\_schedule\_time} - \text{arrival\_time}$ )
- Waiting time: How much thumb-twiddling?
  - Time on the ready queue (not running)

# Performance terminology

- Throughput: How many jobs over time?
  - The rate at which jobs are completed
- Overhead: How much useless work?
  - Time lost due to switching between jobs
- Fairness: How equally are jobs treated?
  - Jobs get same amount of CPU time over some interval

# Workload assumptions (to be relaxed later)

- ① Each job runs for the same amount of time
- ② All jobs arrive at the same time
- ③ All jobs only use the CPU (no I/O)
- ④ Run-time of each job is known

# Scheduling basics

## Workloads:

arrival\_time  
run\_time

## Scheduling

### Policies:

FIFO  
SJF (SJN, SPN)  
STCF  
RR

## Metrics:

turnaround\_time  
response\_time



# FIFO

- Run jobs to completion in arrival\_time order
  - aka FCFS (first come first served)
- simple, minimal context switch overhead

JOB	arrival_time (s)	run_time (s)
A	~0	10
B	~0	10
C	~0	10


## Time

0 A arrives  
0 B arrives  
0 C arrives  
0 run A  
10 complete A  
10 run B  
20 complete B  
20 run C  
30 complete C

# FIFO: Identical jobs

- **Gantt chart**: illustrates how jobs are scheduled over time on a CPU

JOB	arrival_time (s)	run_time (s)
A	~0	10
B	~0	10
C	~0	10



- What is the average turnaround time?
  - $(10+20+30)/3 = 20s$

# Scheduling basics

## Workloads:

arrival\_time  
run\_time

## Scheduling

### Policies:

FIFO  
SJF (SJN, SPN)  
STCF  
RR

## Metrics:

turnaround\_time  
response\_time


# Workload assumptions

- ~~1. Each job runs for the same amount of time~~
2. All jobs arrive at the same time
3. All jobs only use the CPU (no I/O)
4. Run-time of each job is known

## FIFO: Convoy effect

- Problem: turnaround time can suffer when short jobs must wait for long jobs

JOB	arrival_time (s)	run_time (s)
A	~0	60
B	~0	10
C	~0	10



Average turnaround time: 70s


# Shortest Job First (SJF)

- Idea: choose job with the smallest run\_time
  - aka shortest job next (SJN), shortest process next (SPN)
- Consider our previous example

JOB	arrival_time (s)	run_time (s)
A	~0	60
B	~0	10
C	~0	10

What is the average turnaround time with SJF?

# SJF turnaround time



- What is the average turnaround time?
  - $(80+10+20)/3 = 36.7s$
- Average turnaround with FIFO: 70s

# Scheduling basics

## Workloads:

arrival\_time  
run\_time

## Scheduling

### Policies:

FIFO  
SJF (SJN, SPN)  
STCF  
RR

## Metrics:


turnaround\_time  
response\_time



# Workload assumptions

- ~~1. Each job runs for the same amount of time~~
- ~~2. All jobs arrive at the same time~~
3. All jobs only use the CPU (no I/O)
4. Run-time of each job is known

# SJF with late arrivals



JOB	arrival_time (s)	run_time (s)
A	~0	60
B	~10	10
C	~10	10

- What is the average turnaround time?
  - $(60 + (70-10) + (80-10)) / 3 = 63.3s$

# Preemptive scheduling

- FIFO and SJF are non-preemptive
  - Only schedule new job when previous job voluntarily relinquishes CPU
- Preemptive: potentially schedule different job at any point by taking CPU away from running job



# Shortest Time-to-Completion First (STCF)

- SJF + preemption
- Select job with the least remaining time to run next
- Consider our previous example

JOB	arrival_time (s)	run_time (s)
A	~0	60
B	~10	10
C	~10	10

- Average turnaround time with STCF?

# SJF vs. STCF



- Average turnaround time with STCF
  - $((80-0) + (20-10) + (30-10))/3 = 36.7s$
- Average turnaround time with SJF
  - 63.3s

A preemptive scheduler is best characterized by its ability to carry out this state transition:

- 1 READY  $\rightarrow$  RUNNING
- 2 RUNNING  $\rightarrow$  READY

Which of the following state transitions will never take place in a FIFO scheduler?

- ① READY  $\rightarrow$  RUNNING
- ② RUNNING  $\rightarrow$  READY

# Scheduling basics

## Workloads:

arrival\_time  
run\_time

## Scheduling

### Policies:

FIFO  
SJF (SJN, SPN)  
STCF  
RR

## Metrics:

turnaround\_time  
response\_time



# Response vs. turnaround

- Response time: first run time - arrival time

B's turnaround: 20s  $\longleftrightarrow$

B's response: 10s  $\longleftrightarrow$




- FIFO, SJF, and STCF can have poor response time


# Round-Robin (RR)

- Each job allowed to run for a quantum
  - quantum = some configured period of time
- Context is switched (at the latest) at the end of the quantum – preemption!
- Next job is the one on the ready queue that hasn't run for the longest amount of time

# FIFO vs RR



Avg Response Time?  
 $(0+5+10)/3 = 5$



Avg Response Time?  
 $(0+1+2)/3 = 1$

- In what way is RR worse?
  - Avg. turnaround time with equal job length is horrible
- Other reasons why RR could be better?
  - If don't know run-time of each job, gives short jobs a chance to run and finish fast


# Time quantum

- What is a good quantum size?
- Too long, and it morphs into FIFO
- Too short, and time is wasted on context switching
- Typical quantum: about 100X cost of context switch ( $\sim 100$  ms vs.  $\ll 1$  ms)

# RR Example

- time slice = 20
- workload (with runtimes)
  - P1: 53
  - P2: 8
  - P3: 68
  - P4: 24

# RR Example



Avg turnaround time =  $(125+28+153+112)/4 = 104.25$

# Scheduling basics

## Workloads:

arrival\_time  
run\_time

## Scheduling

### Policies:

FIFO  
SJF (SJN, SPN)  
STCF  
RR

## Metrics:

turnaround\_time  
response\_time



# Workload assumptions

- ~~1. Each job runs for the same amount of time~~
- ~~2. All jobs arrive at the same time~~
- ~~3. All jobs only use the CPU (no I/O)~~
4. Run-time of each job is known

# Incorporating I/O

- When a job initiates an I/O request
  - The job is blocked waiting for I/O completion
  - The scheduler should schedule another job on the CPU
- When the I/O completes
  - An interrupt is raised
  - The OS moves the process from blocked back to the ready state

# I/O bound and CPU bound jobs

- A job is CPU bound if it needs lots of CPU time to progress
  - These jobs have a lot of logic and math

# I/O bound and CPU bound jobs

- A job is CPU bound if it needs lots of CPU time to progress
  - These jobs have a lot of logic and math
  - Usually in running or ready state

# I/O bound and CPU bound jobs

- A job is CPU bound if it needs lots of CPU time to progress
  - These jobs have a lot of logic and math
  - Usually in running or ready state
- Job is I/O bound if it needs to do lots of I/O to progress
  - These jobs access disk, network, etc

# I/O bound and CPU bound jobs

- A job is CPU bound if it needs lots of CPU time to progress
  - These jobs have a lot of logic and math
  - Usually in running or ready state
- Job is I/O bound if it needs to do lots of I/O to progress
  - These jobs access disk, network, etc
  - Usually in the waiting (or, blocked) state

- Consider a mixture of I/O-bound and CPU-bound jobs
- Even though the I/O completes quickly, the I/O-bound task must wait to be reassigned the CPU until the CPU-bound tasks both complete their time slice

