

Operating Systems

Autumn 2024

race conditions, mutual exclusion

What is a thread?

- An abstraction for a single running process
- A multi-threaded program has more than one point of execution
 - Multiple program counters, one for each thread
 - They share code, heap, and global variables
 - Each thread has its own stack
 - Each thread has its own private set of registers

Data race

- A program has a data race if it is possible for a thread to modify an addressable location at the same time that another thread is accessing the same location
- The result/correctness depends on the sequence/timing of events; i.e., how things end up being scheduled.

Shared counter

- Two threads updating a single shared variable `cnt`

```
int cnt = 0;
void *worker(void *ptr) {
    int i;
    for (i = 0; i < ITERATIONS_PER_THREAD; i++)
        cnt++;
}
```

- What happens when two threads execute concurrently?

Order of execution of threads

- Possible implementation of `cnt++`:

```
eax := cnt      // read
eax := eax + 1   // modify
cnt := eax       // write
```

- OS might decide to context switch from one thread to another at any time
- Thus the atomic actions of concurrent threads may be interleaved in any possible order

Possible interleaving

- `eax_i` denotes the value of register in thread `i`



```
eax_1 := cnt
eax_1 := eax_1 + 1
<switch>
eax_2 := cnt
eax_2 := eax_2 + 1
cnt := eax_2
<switch>
cnt := eax_1
```

- Result: `cnt` one less than correct! (lost update)

- Consider 2 threads sharing a global variable `count`, initially 10:

```
// Thread A  
count++;  
  
// Thread B  
count--;
```

- What are the possible values for `count` after both threads finish executing?

```
int count = 0;
// Thread 1
for (i=0; i < 10; i++) {
    count++;
}

// Thread 2
for (i=0; i < 10; i++) {
    count++;
}
```

- What is the final value of count?
 - A value between 2 and 20

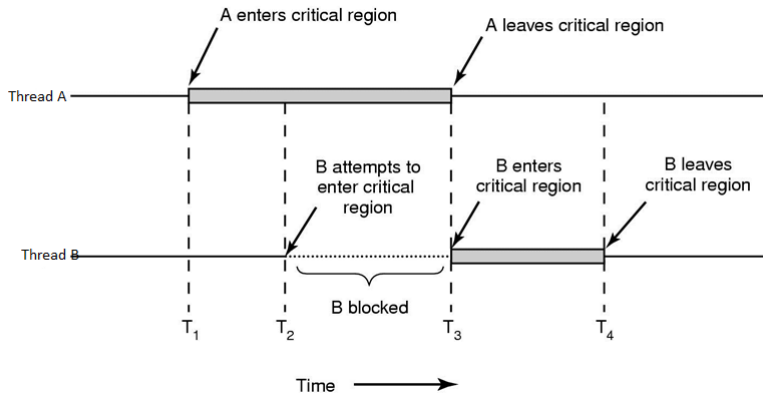
count = 2

- t1 and t2 read count
- t1 increments 9 times. so count becomes 9
- t2 increments once. so count becomes 1
- t1 reads count (value is 1)
- t2 increments 9 times. so count becomes 10
- t1 increments once. so count becomes 2

Critical section

- Critical section: code that access shared resource (e.g., variable or data structure)
- Race condition: arises when multiple threads simultaneously enter critical section leading to a non-deterministic outcome

Critical section



What basic mechanism can stop thread B from entering critical region when thread A is in?

Requirements for a good solution

- Mutual exclusion (safety property)
 - At most one thread can be in CS at a time

Requirements for a good solution

- **Mutual exclusion** (safety property)
 - At most one thread can be in CS at a time
- **Progress** (liveness property)
 - If no thread is currently in CS and threads are trying to access, one should eventually be able to enter the CS

Requirements for a good solution

- Mutual exclusion (safety property)
 - At most one thread can be in CS at a time
- Progress (liveness property)
 - If no thread is currently in CS and threads are trying to access, one should eventually be able to enter the CS
- Bounded waiting (liveness property)
 - Once a thread T starts trying to enter the CS, there is a bound on the number of times other threads get in

Progress vs. Bounded waiting

- Liveness requirements are needed for a solution to be useful

Progress vs. Bounded waiting

- Liveness requirements are needed for a solution to be useful
- Progress
 - If no thread can enter CS, we don't have progress

Progress vs. Bounded waiting

- Liveness requirements are needed for a solution to be useful
- Progress
 - If no thread can enter CS, we don't have progress
- Bounded waiting
 - If thread A is waiting to enter CS while B repeatedly leaves and re-enters CS ad infinitum, we don't have bounded waiting

Mutual exclusion: legacy solutions

- Mutual exclusion algorithm solely based on read and write operations to a shared memory

Mutual exclusion: legacy solutions

- Mutual exclusion algorithm solely based on read and write operations to a shared memory
- First correct solution for two threads by Dekker in 1966
- Peterson proposed a simpler solution in 1981

Mutual exclusion: legacy solutions

Solution for 2 threads T_0 and T_1

Algorithm 1 Peterson's algorithm for thread T_i

Global Variables:

```
1: bool wants[2] = {false, false};  
2: int not_turn; /* can be 0 or 1 */
```

```
3: enter_CS()
```

```
4:   wants[i] = true;
```

```
5:   not_turn = i;
```

```
6:   while wants[1-i] == true and not_turn == i do
```

```
7:     /* do nothing */
```

```
8:   end while
```

```
9: leave_CS()
```

```
10:  wants[i] = false;
```

Peterson's algorithm: a few comments

- wants: To declare that the thread wants to enter
- not_turn: To arbitrate if the 2 threads want to enter
- Line 6: “The other thread wants to access and not our turn, so loop”

Correctness (1)

Algorithm 1 Peterson's algorithm for thread T_i

Global Variables:

```
1: bool wants[2] = {false, false};
2: int not_turn; /* can be 0 or 1 */

3: enter_CS()
4:   wants[i] = true;
5:   not_turn = i;
6:   while wants[1-i] == true and not_turn == i do
7:     /* do nothing */
8:   end while

9: leave_CS()
10:  wants[i] = false;
```

- Mutual exclusion

- Assume both threads want to enter CS, then $wants[0] = wants[1] = 1$
- not_turn can only be 0 or 1
- the condition of the while loop can only resolve to false for one of the threads

Correctness (2)

Algorithm 1 Peterson's algorithm for thread T_i

Global Variables:

```
1: bool wants[2] = {false, false};
2: int not_turn; /* can be 0 or 1 */

3: enter_CS()
4:   wants[i] = true;
5:   not_turn = i;
6:   while wants[1-i] == true and not_turn == i do
7:     /* do nothing */
8:   end while

9: leave_CS()
10:  wants[i] = false;
```

- Progress

- If T_0 doesn't want CS, $\text{wants}[0] == \text{false}$, so T_1 won't loop
- If both threads try to enter, one thread will succeed

Correctness (3)

Algorithm 1 Peterson's algorithm for thread T_i

Global Variables:

```
1: bool wants[2] = {false, false};
2: int not_turn; /* can be 0 or 1 */

3: enter_CS()
4:   wants[i] = true;
5:   not_turn = i;
6:   while wants[1-i] == true and not_turn == i do
7:     /* do nothing */
8:   end while

9: leave_CS()
10:  wants[i] = false;
```

- **Bounded waiting**

- If T_1 was blocked from CS and T_0 tries to re-enter, T_0 will set not_turn = 0, allowing T_1 in

- Consider a modified algorithm where lines 4 and 5 are swapped
- Does this solution work?

Global Variables:

```
1: bool wants[2] = {false, false};  
2: int not_turn; /* can be 0 or 1 */
```

```
3: enter CS()
```

```
4:   not_turn = i;
```

```
5:   wants[i] = true;
```

```
6:   while wants[1-i] == true and not_turn == i do
```

```
7:     /* do nothing */
```

```
8:   end while
```

```
9: leave CS()
```

```
10:  wants[i] = false;
```

- Mutual exclusion violation
- Bad trace
 - T1 does `not_turn = 1`
 - T0 does `not_turn = 0`
 - T0 does `wants[0] = true`
 - T0 does `while (wants[1] == true and not_turn == 0); // (false and true) → false`

- Mutual exclusion violation
- Bad trace
 - T1 does `not_turn = 1`
 - T0 does `not_turn = 0`
 - T0 does `wants[0] = true`
 - T0 does `while (wants[1] == true and not_turn == 0); // (false and true) → false`
 - T0 enters the CS

- Mutual exclusion violation
- Bad trace
 - T1 does `not_turn = 1`
 - T0 does `not_turn = 0`
 - T0 does `wants[0] = true`
 - T0 does `while (wants[1] == true and not_turn == 0); // (false and true) → false`
 - T0 enters the CS
 - T1 does `wants[1] = true`
 - T1 does `while (wants[0] == true and not_turn == 1); // (true and false) → false`

- Mutual exclusion violation
- Bad trace
 - T1 does `not_turn = 1`
 - T0 does `not_turn = 0`
 - T0 does `wants[0] = true`
 - T0 does `while (wants[1] == true and not_turn == 0); // (false and true) → false`
 - T0 enters the CS
 - T1 does `wants[1] = true`
 - T1 does `while (wants[0] == true and not_turn == 1); // (true and false) → false`
 - T1 enters the CS

- Assume the `while` statement in line 6 is changed to
 - `while (wants[1-i] == true and not_turn != i)`
- Show a trace in which there is a mutual exclusion violation
- Show a trace in which there is NO mutual exclusion violation

Peterson's algorithm: a few more comments

- Given solution works for 2 threads
- Can be generalized to n threads but n must be known in advance
- To implement a general lock, processors provide hardware primitives