

Operating Systems

Autumn 2024

sleeping locks

Recap: Thread states

- ready: waiting to be assigned to CPU
 - could run, but another thread has the CPU

Recap: Thread states

- ready: waiting to be assigned to CPU
 - could run, but another thread has the CPU
- running: executing on the CPU
 - is the thread that currently controls the CPU

Recap: Thread states

- ready: waiting to be assigned to CPU
 - could run, but another thread has the CPU
- running: executing on the CPU
 - is the thread that currently controls the CPU
- waiting: waiting for an event, e.g. I/O
 - cannot make progress until event happens

Spinlock performance

- When context switching occurs inside a CS, “other” threads start to spin endlessly, waiting for the interrupted (lock-holding) thread to be run again

Spinlock performance

- When context switching occurs inside a CS, “other” threads start to spin endlessly, waiting for the interrupted (lock-holding) thread to be run again
- Instead of spinning for a lock, a contending thread could simply give up the CPU and check back later

Spinlock performance

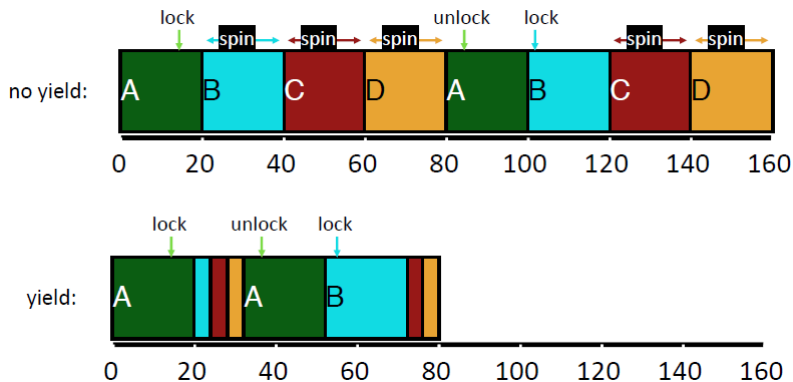
- When context switching occurs inside a CS, “other” threads start to spin endlessly, waiting for the interrupted (lock-holding) thread to be run again
- Instead of spinning for a lock, a contending thread could simply give up the CPU and check back later
 - `yield` system call moves calling thread from running to ready state

Spinlock performance

- When context switching occurs inside a CS, “other” threads start to spin endlessly, waiting for the interrupted (lock-holding) thread to be run again
- Instead of spinning for a lock, a contending thread could simply give up the CPU and check back later
 - **yield** system call moves calling thread from running to ready state

```
void lock(lock_t *lock) {  
    while (xchg(&lock->flag, 1) == 1)  
        yield(); // give up the CPU  
}
```


Why is yield() useful?




Context switching

- How about 100 threads on one CPU?
- If one thread acquires the lock and is preempted before releasing it, the other 99 will call `lock()`, find the lock held, and yield the CPU (99 context switches)

Context switching

- How about 100 threads on one CPU?
- If one thread acquires the lock and is preempted before releasing it, the other 99 will call lock(), find the lock held, and yield the CPU (99 context switches)
 - better than spinlock which wastes 99 time slices spinning

Context switching

- 
- How about 100 threads on one CPU?
 - If one thread acquires the lock and is preempted before releasing it, the other 99 will call `lock()`, find the lock held, and yield the CPU (99 context switches)
 - better than spinlock which wastes 99 time slices spinning
 - Even with yield, high context switch cost

Blocking locks (aka mutexes): sleeping instead of spinning

- recall that scheduler only runs ready threads

Blocking locks (aka mutexes): sleeping instead of spinning

- recall that scheduler only runs ready threads
- idea: remove waiting threads from the scheduler's ready queue and put them on waiting queue
- no time wasted on threads that are contending on the lock

Blocking locks (aka mutexes): sleeping instead of spinning

- recall that scheduler only runs ready threads
- idea: remove waiting threads from the scheduler's ready queue and put them on waiting queue
- no time wasted on threads that are contending on the lock
- when lock is released, one thread on queue is restarted

Blocking locks: implementation idea

- 1 wait queue and 1 lock variable per blocking lock
- Implementing blocking lock gets tricky
 - Need to enforce mutual exclusion while performing operations on the wait queue (adding/removing) and the lock variable

Blocking locks: implementation idea

- 1 wait queue and 1 lock variable per blocking lock
- Implementing blocking lock gets tricky
 - Need to enforce mutual exclusion while performing operations on the wait queue (adding/removing) and the lock variable
 - how to implement this?

Blocking locks: implementation idea

- 1 wait queue and 1 lock variable per blocking lock
- Implementing blocking lock gets tricky
 - Need to enforce mutual exclusion while performing operations on the wait queue (adding/removing) and the lock variable
 - how to implement this?
 - use spinlock in the implementation of blocking lock!

Blocking locks: implementation idea

- Two separate levels of locking
 - holding spinlock guarding queue/variable from concurrent modification
 - holding actual blocking lock

Blocking locks: data structure

```
typedef struct {  
    int flag;  
    queue_t *q;  
    int guard;  
} lock_t;
```

tracks whether any thread has locked
and not unlocked

Blocking locks: data structure

```
typedef struct {  
    int flag;  
    queue_t *q;  
    int guard;  
} lock_t;
```

list of threads that discovered lock is taken and are waiting for it to be free. these threads are not in scheduler's ready queue

Blocking locks: data structure

```
typedef struct {  
    int flag;  
    queue_t *q;  
    int guard;  
} lock_t;
```

used as a spinlock to protect flag and q manipulation

Blocking locks: OS support

- `park()`
 - Put a calling thread to sleep (remove it from runqueue and ready queue)

Blocking locks: OS support

- `park()`
 - Put a calling thread to sleep (remove it from runqueue and ready queue)
- `unpark(threadID)`
 - Wake a particular thread designated by `threadID`

Blocking locks: OS support

- `park()`
 - Put a calling thread to sleep (remove it from runqueue and ready queue)
- `unpark(threadID)`
 - Wake a particular thread designated by `threadID`

Blocking locks: OS support

- `park()`
 - Put a calling thread to sleep (remove it from runqueue and ready queue)
- `unpark(threadID)`
 - Wake a particular thread designated by `threadID`
- `park`, `unpark` inspired by Solaris

lock() code

```
void lock(lock_t *m) {
    while (xchg(&m->guard, 1) == 1)
        ; // acquire guard lock
    if (m->flag == 0) {
        m->flag = 1; // lock acquired
        m->guard = 0;
    } else {
        queue_add(m->q, getpid());
        m->guard = 0;
        park();
    }
}
```

```
typedef struct {
    int flag;
    int guard;
    queue_t *q;
}
```

- What is protected by the spin lock guard?

lock() code

```
void lock(lock_t *m) {
    while (xchg(&m->guard, 1) == 1)
        ; // acquire guard lock
    if (m->flag == 0) {
        m->flag = 1; // lock acquired
        m->guard = 0;
    } else {
        queue_add(m->q, getpid());
        m->guard = 0;
        park();
    }
}
```

```
typedef struct {
    int flag;
    int guard;
    queue_t *q;
}
```

- What is protected by the spin lock guard?
 - the flag and queue manipulation code
 - no more than a single thread can ever be active within that code

lock() code

```
void lock(lock_t *m) {  
    while (xchg(&m->guard, 1) == 1)  
        ; // acquire guard lock  
    if (m->flag == 0) {  
        m->flag = 1; // lock acquired  
        m->guard = 0;  
    } else {  
        queue_add(m->q, getpid());  
        m->guard = 0;  
        park();  
    }  
}
```

```
typedef struct {  
    int flag;  
    int guard;  
    queue_t *q;  
}
```

- What is protected by the spin lock guard?
 - the flag and queue manipulation code
 - no more than a single thread can ever be active within that code
- Why is this better than a simple spin lock?

lock() code

```
void lock(lock_t *m) {  
    while (xchg(&m->guard, 1) == 1)  
        ; // acquire guard lock  
    if (m->flag == 0) {  
        m->flag = 1; // lock acquired  
        m->guard = 0;  
    } else {  
        queue_add(m->q, gettid());  
        park();  
        m->guard = 0;  
    }  
}
```

```
typedef struct {  
    int flag;  
    int guard;  
    queue_t *q;  
}
```

- Setting `m->guard = 0` after calling `park()`
- Would it work?

unlock() code

```
void unlock(lock_t *m) {  
    while (xchg(&m->guard, 1) == 1)  
        ; // acquire guard lock  
    if (queue_empty(m->q))  
        m->flag = 0; // let go of lock  
    else  
        // hold lock (for next thread!)  
        unpark(queue_remove(m->q));  
    m->guard = 0;  
}
```

```
typedef struct {  
    int flag;  
    int guard;  
    queue_t *q;  
} lock_t;
```

- In unlock, there is no setting of flag=0 when we unpark. Why?

Race-condition bug

Thread 1 in lock

```
if (m->flag) {  
    queue_add(m->q, getpid());  
    m->guard = 0;  
  
    park();  
}
```

Thread 2 in unlock

```
while (xchg(&m->guard, 1) == 1)  
    ;  
if (queue_empty(m->q))  
    m->flag = 0;  
else  
    unpark(queue_remove(m->q));
```

- Assume Thread 2 is holding the lock
- Thread 1 calls lock()
- Before the call to park(), a switch to Thread 2 happens
- Thread 2 calls unlock() and does unpark()
- When Thread 1 calls park(), it sleeps forever (why?)

Solving the race problem: final correct lock

- `setpark()`
 - informs OS of my plan to `park()` myself
- If there is an `unpark()` between my `setpark()` and `park()`, `park()` will return immediately (no blocking)

```
queue_add(m->q, getpid());  
setpark(); | // new code  
m->guard=0;  
park();
```

- park, unpark and setpark inspired by Solaris
- Other OSes provide different mechanisms to support blocking synchronization
- E.g., Linux has a mechanism called **futex**
 - it renders guard and setpark unnecessary
- Read more about futex in OSTEP (brief)