

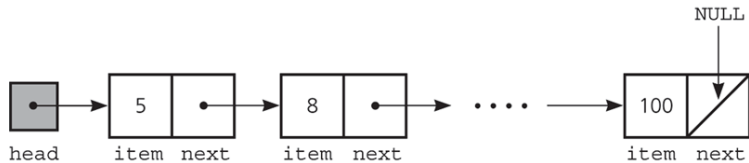
# IE 411: Operating Systems

Linked lists: locking and lock-free

# Concurrent data structures

- For a data structure, we would like multiple local (independent) operations to be allowed concurrently
- Can use locks to achieve thread-safe access
- But let's see if we can do thread-safe access without any locks at all
- We will illustrate the main ideas using linked lists

# A sorted linked list

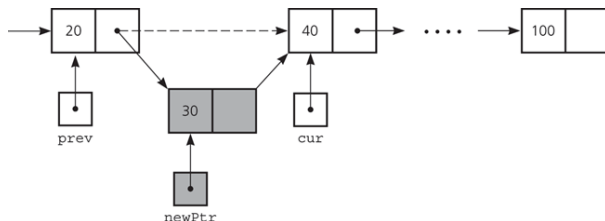


# A sorted linked list

- To insert a node search for the insertion point (*cur*, *prev*)

`newPtr->next = cur;`

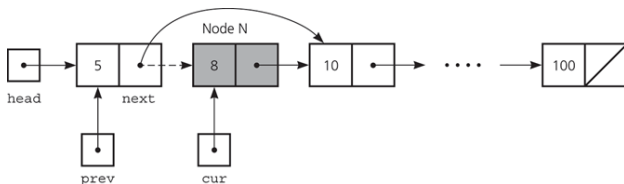
`prev->next = newPtr;`



# A sorted linked list

- To delete search for node

$prev \rightarrow next = cur \rightarrow next;$

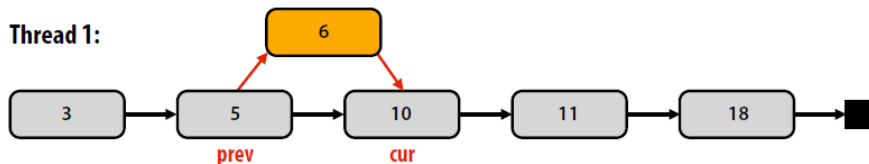


# Race condition

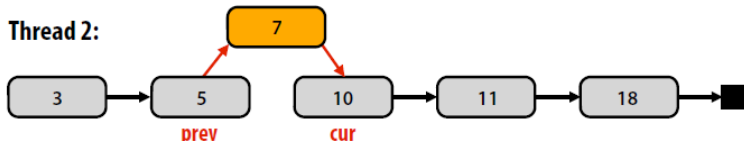
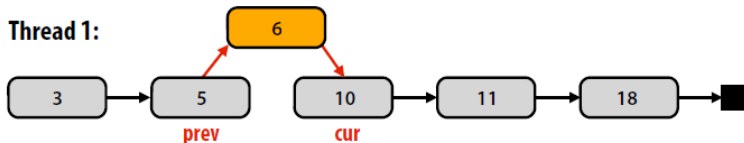
- Thread 1 attempts to insert 6
- Thread 2 attempts to insert 7



**Thread 1:**

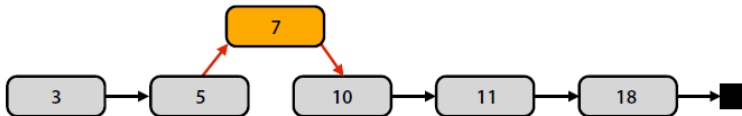


# Race condition



Thread 1 and thread 2 both compute same prev and cur.  
Result: one of the insertions gets lost!

**Result: (assuming thread 1 updates prev->next before thread 2)**



# Single global lock

- Use a per-list lock
- Advantages
  - simple to implement



# Single global lock

- Use a per-list lock
- Advantages
  - simple to implement
- Disadvantages?

# Single global lock

- Use a per-list lock
- Advantages
  - simple to implement
- Disadvantages?
  - Operations on the data structure are serialized

# Single global lock

- Use a per-list lock
- Advantages
  - simple to implement
- Disadvantages?
  - Operations on the data structure are serialized
  - May limit application performance

# Lock-free algorithms

- protecting DS (e.g. BST, linked list) with a single lock is pessimistic as it assumes conflicts will occur
- a lockless algorithm is optimistic as it assumes conflicts unlikely to occur and, when they are detected, they are resolved

# Lock-free algorithms

- protecting DS (e.g. BST, linked list) with a single lock is pessimistic as it assumes conflicts will occur
- a lockless algorithm is optimistic as it assumes conflicts unlikely to occur and, when they are detected, they are resolved
- Advantages compared to locking?

# Lock-free algorithms

- protecting DS (e.g. BST, linked list) with a single lock is pessimistic as it assumes conflicts will occur
- a lockless algorithm is optimistic as it assumes conflicts unlikely to occur and, when they are detected, they are resolved
- Advantages compared to locking?
  - allows concurrency while there are no conflicts **which hopefully is so most of the time**

# Atomic Compare-and-Swap (CAS)

```
bool CAS(  
    memory location L,  
    expected value V at L,  
    desired new value V1 at L  
);
```

If (the expected value V at memory location L == the current value at L), CAS succeeds by storing the the desired value V1 at L and returns TRUE.

# Adding nodes

- One list, three nodes: a, d, e
- To insert node c

① `c->next = d`

② ATOMICALLY

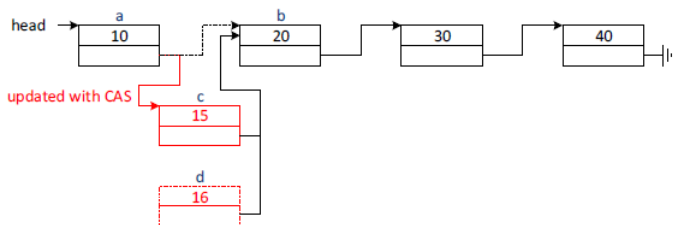
```
if (a->next == d)
    a->next = c
else
    fail
```

- This translates into `CAS(&a->next, d, c)`
- CAS succeeded: c was successfully inserted between a and d
- CAS failed: retry



# Adding nodes

- if 2 threads try to add nodes at the same position

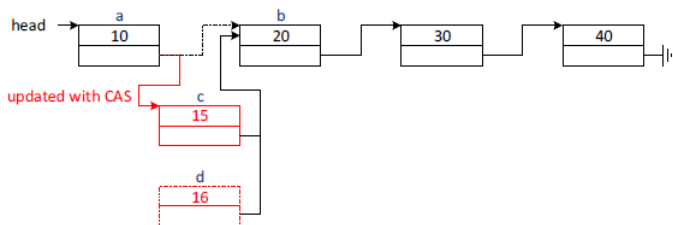


```
CAS(&a->next, b, c); // first CAS executed will succeed..  
CAS(&a->next, b, d); // and thus second CAS executed will FAIL
```

- first CAS executed succeeds, second will fail as `a->next != b`

# Adding nodes

- if 2 threads try to add nodes at the same position

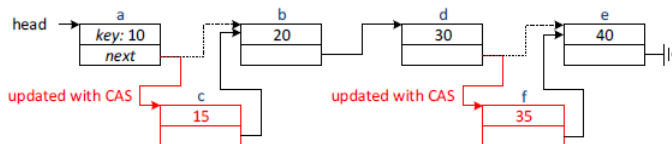


```
CAS(&a->next, b, c); // first CAS executed will succeed..  
CAS(&a->next, b, d); // and thus second CAS executed will FAIL
```

- first CAS executed succeeds, second will fail as `a->next != b`
- RETRY on failure, which means searching for insertion point AGAIN and, if key not found, set up and re-execute CAS

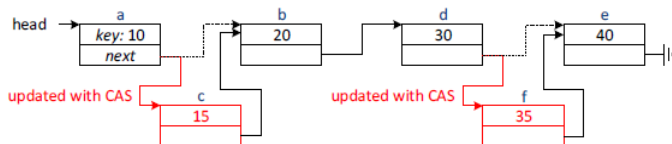
# Adding nodes

- Ex: use CAS to add nodes 15 and 35



# Adding nodes

- Ex: use CAS to add nodes 15 and 35



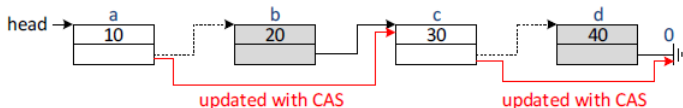
- search for insertion point, initialise next pointer and then execute with correct parameters to insert node into list

```
CAS(&a->next, b, c);    // add node c between a and b
CAS(&d->next, e, f);    // add node f between d and e
```

- disjoint-access parallelism

# Removing nodes

- search for node and then execute CAS with correct parameters to remove node from list
- consider 2 threads removing non-adjacent nodes



```
CAS(&a->next, b, c); // remove node b (20)
```

```
CAS(&c->next, d, 0); // remove node d (40)
```

- disjoint access parallelism

# Removing nodes

- if two threads try to remove the same node



```
CAS(&a->next, b, c);
```

```
CAS(&a->next, b, c);
```

# Removing nodes

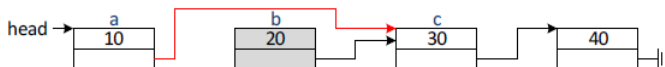
- if two threads try to remove the same node



```
CAS(&a->next, b, c);
```

```
CAS(&a->next, b, c);
```

- assume first CAS executed succeeds



- then second CAS executed fails as  $a \rightarrow next \neq b$

# Removing nodes

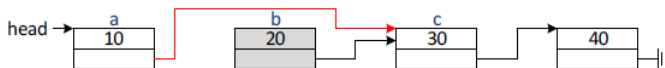
- if two threads try to remove the same node



```
CAS(&a->next, b, c);
```

```
CAS(&a->next, b, c);
```

- assume first CAS executed succeeds

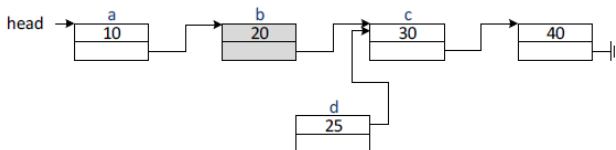


- then second CAS executed fails as  $a \rightarrow \text{next} \neq b$
- RETRY on failure, which means searching AGAIN for node (which may not be found)

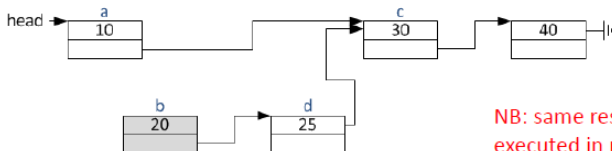


# What doesn't work...

- consider removing node 20 and adding node 25 concurrently



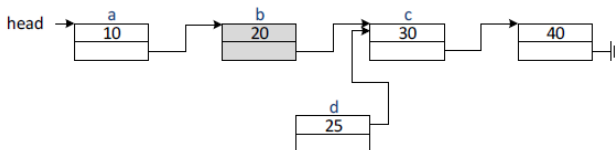
```
CAS(&a->next, b, c); // remove 20  
CAS(&b->next, c, d); // add 25
```



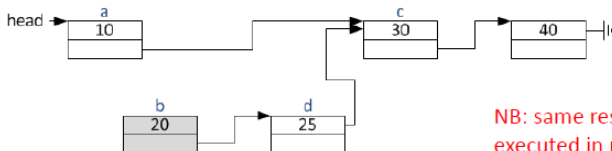
NB: same result if CAS instructions  
executed in reverse order

# What doesn't work...

- consider removing node 20 and adding node 25 concurrently



```
CAS(&a->next, b, c); // remove 20  
CAS(&b->next, c, d); // add 25
```

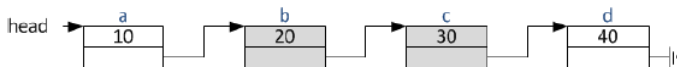


NB: same result if CAS instructions  
executed in reverse order

- NOT what was intended!

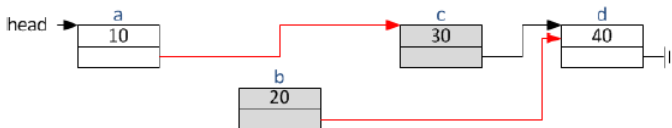
# What doesn't work...

- imagine deleting adjacent nodes



```
CAS(&a->next, b, c); // remove 20
```

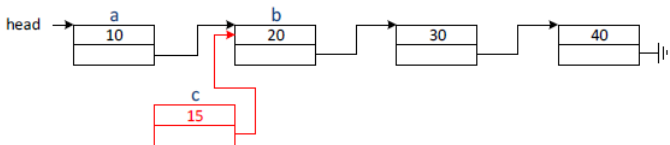
```
CAS(&b->next, c, d); // remove 30
```



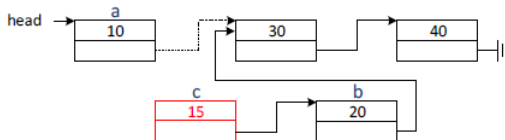
- AGAIN NOT what was intended!

# ABA Problem

- imagine insertion point found, BUT before `CAS(&a->next, b, c)` is executed, thread is pre-empted

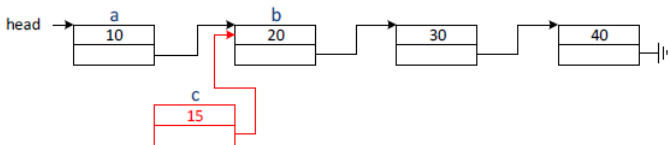


- another thread then removes b from list

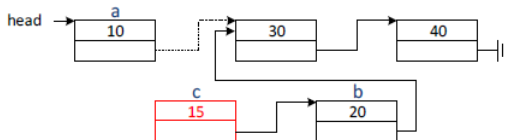


# ABA Problem

- imagine insertion point found, BUT before `CAS(&a->next, b, c)` is executed, thread is pre-empted



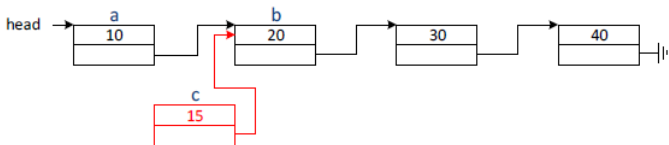
- another thread then removes b from list



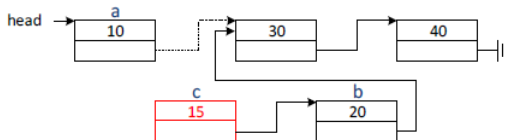
- if thread adding 15 resumes execution, the CAS fails which is OK in this case

# ABA Problem

- imagine insertion point found, BUT before `CAS(&a->next, b, c)` is executed, thread is pre-empted



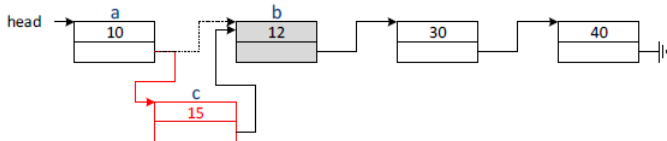
- another thread then removes b from list



- if thread adding 15 resumes execution, the CAS fails which is OK in this case
- BUT what bad thing can happen?

# ABA Problem

- if the memory used by b is reused, for example by a thread adding key 12 to the list before thread adding 15 resumes ...
- when the thread adding 15 to list resumes, its CAS will succeed and 15 will be added into the list at the wrong position



# ABA Problem

- avoid the ABA problem by not reusing nodes:
  - nodes cannot be reused if any thread has or can get a pointer to the node



# ABA Problem

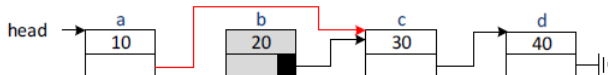
- avoid the ABA problem by not reusing nodes:
  - nodes cannot be reused if any thread has or can get a pointer to the node
- Disadvantages?

# ABA Problem

- avoid the ABA problem by not reusing nodes:
  - nodes cannot be reused if any thread has or can get a pointer to the node
- Disadvantages?
  - will quickly run out of memory

# Marked nodes

- Use two step removal e.g. `remove(20)`



- 1 atomically mark node by setting LSB of next pointer (logically removes node)
- 2 remove node by updating next pointer using CAS

# Marked nodes

- Marked node indicated by an ODD address in its next field
  - OK as addresses normally aligned on at least 4 byte boundary [2 or 3 LSBs normally 0]
- e.g., to atomically mark node b [logically remove]



# Marked nodes

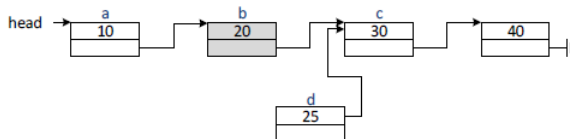
- Marked node indicated by an ODD address in its next field
  - OK as addresses normally aligned on at least 4 byte boundary [2 or 3 LSBs normally 0]
- e.g., to atomically mark node b [logically remove]



```
CAS(&b->next, c, c+1) //assumes node UNMARKED
```

# Revisit: adding node and removing node

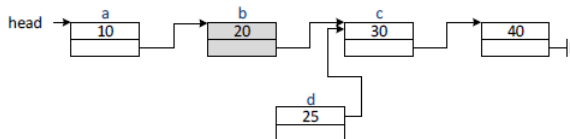
- imagine adding node [25] and removing node [20] concurrently



- (1) `CAS(&b->next, c, d);`    // add 25  
    and
- (2) `if (CAS(&b->next, c, c+1) == 1)` // MARK node b and then
- (3)     `CAS(&a->next, b, c);`            // remove b [20]

# Revisit: adding node and removing node

- imagine adding node [25] and removing node [20] concurrently



(1) `CAS(&b->next, c, d);`    `// add 25`

and

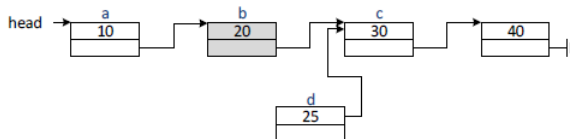
(2) `if (CAS(&b->next, c, c+1) == 1)` `// MARK node b and then`

(3)        `CAS(&a->next, b, c);`                `// remove b [20]`

- if (1) executed first, (2) will fail as `b->next != c`

# Revisit: adding node and removing node

- imagine adding node [25] and removing node [20] concurrently



(1) `CAS(&b->next, c, d);`    // add 25

and

(2) `if (CAS(&b->next, c, c+1) == 1)` // MARK node b and then

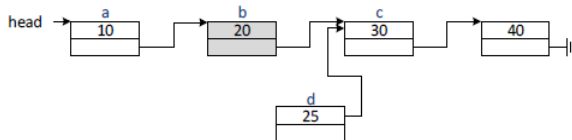
(3)        `CAS(&a->next, b, c);`                // remove b [20]

- if (1) executed first, (2) will fail as `b->next != c`
- if (2) executed first, (1) will fail as `b->next != c`



# Revisit: adding node and removing node

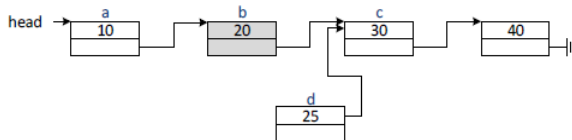
- imagine adding node [25] and removing node [20] concurrently



- (1) `CAS(&b->next, c, d);`    `// add 25`  
and
- (2) `if (CAS(&b->next, c, c+1) == 1)` `// MARK node b and then`
- (3)        `CAS(&a->next, b, c);`            `// remove b [20]`

# Revisit: adding node and removing node

- imagine adding node [25] and removing node [20] concurrently



- (1) `CAS(&b->next, c, d);`    // add 25  
    and
  - (2) `if (CAS(&b->next, c, c+1) == 1)` // MARK node b and then
  - (3)     `CAS(&a->next, b, c);`            // remove b [20]
- if (3) fails, it means that **a** no longer points to **b**, BUT b is logically marked and can be removed later
    - OK for list to contain temporary marked nodes

# What still needs to be done?

- Previous solution avoids ABA problem by NOT re-using nodes
  - there is no code for freeing or reusing nodes

# What still needs to be done?

- Previous solution avoids ABA problem by NOT re-using nodes
  - there is no code for freeing or reusing nodes
- Solutions with memory management:
  - A Pragmatic Implementation of Non-Blocking Linked Lists, Tim Harris, 2001
  - Hazard Pointers: Safe Memory Reclamation for Lock-Free Objects, Maged M. Michael, 2004